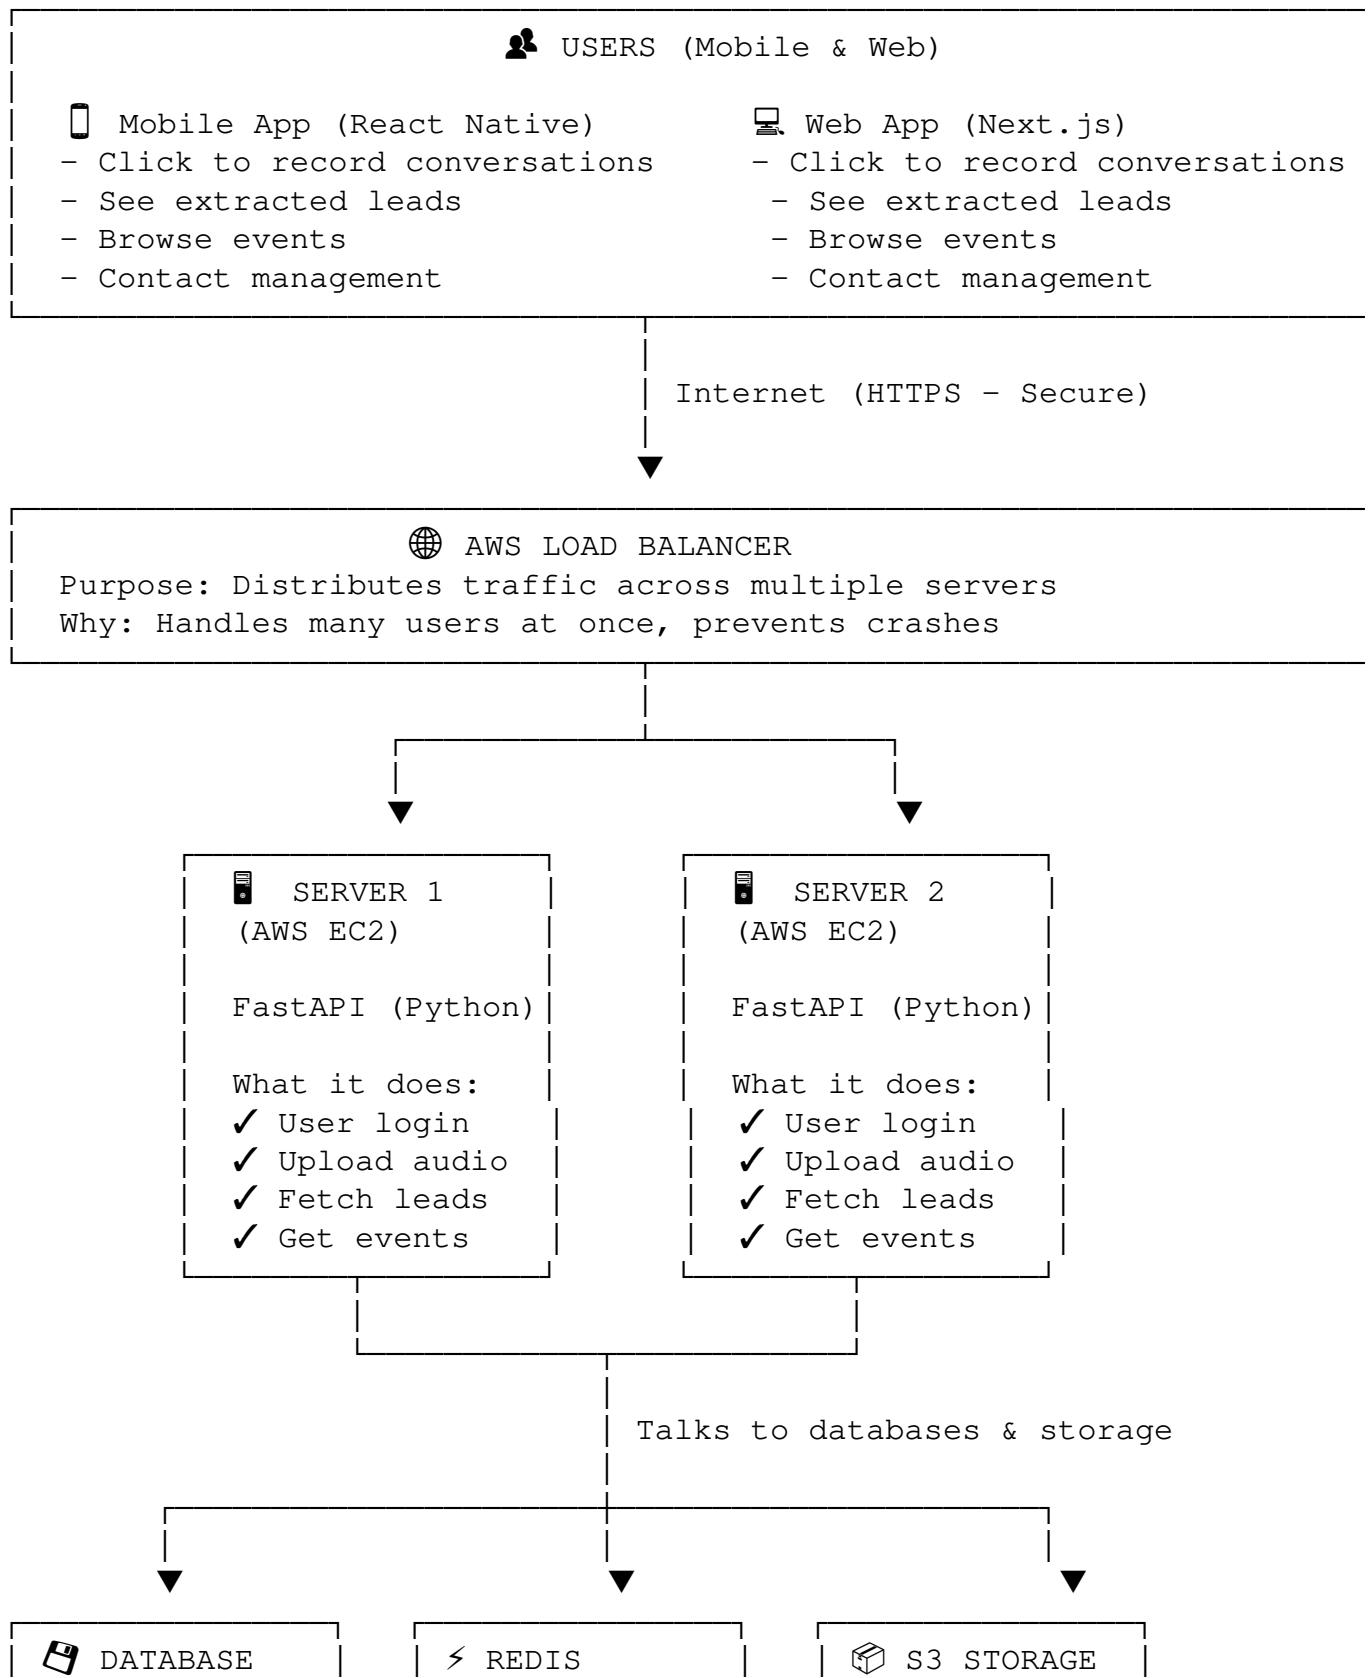
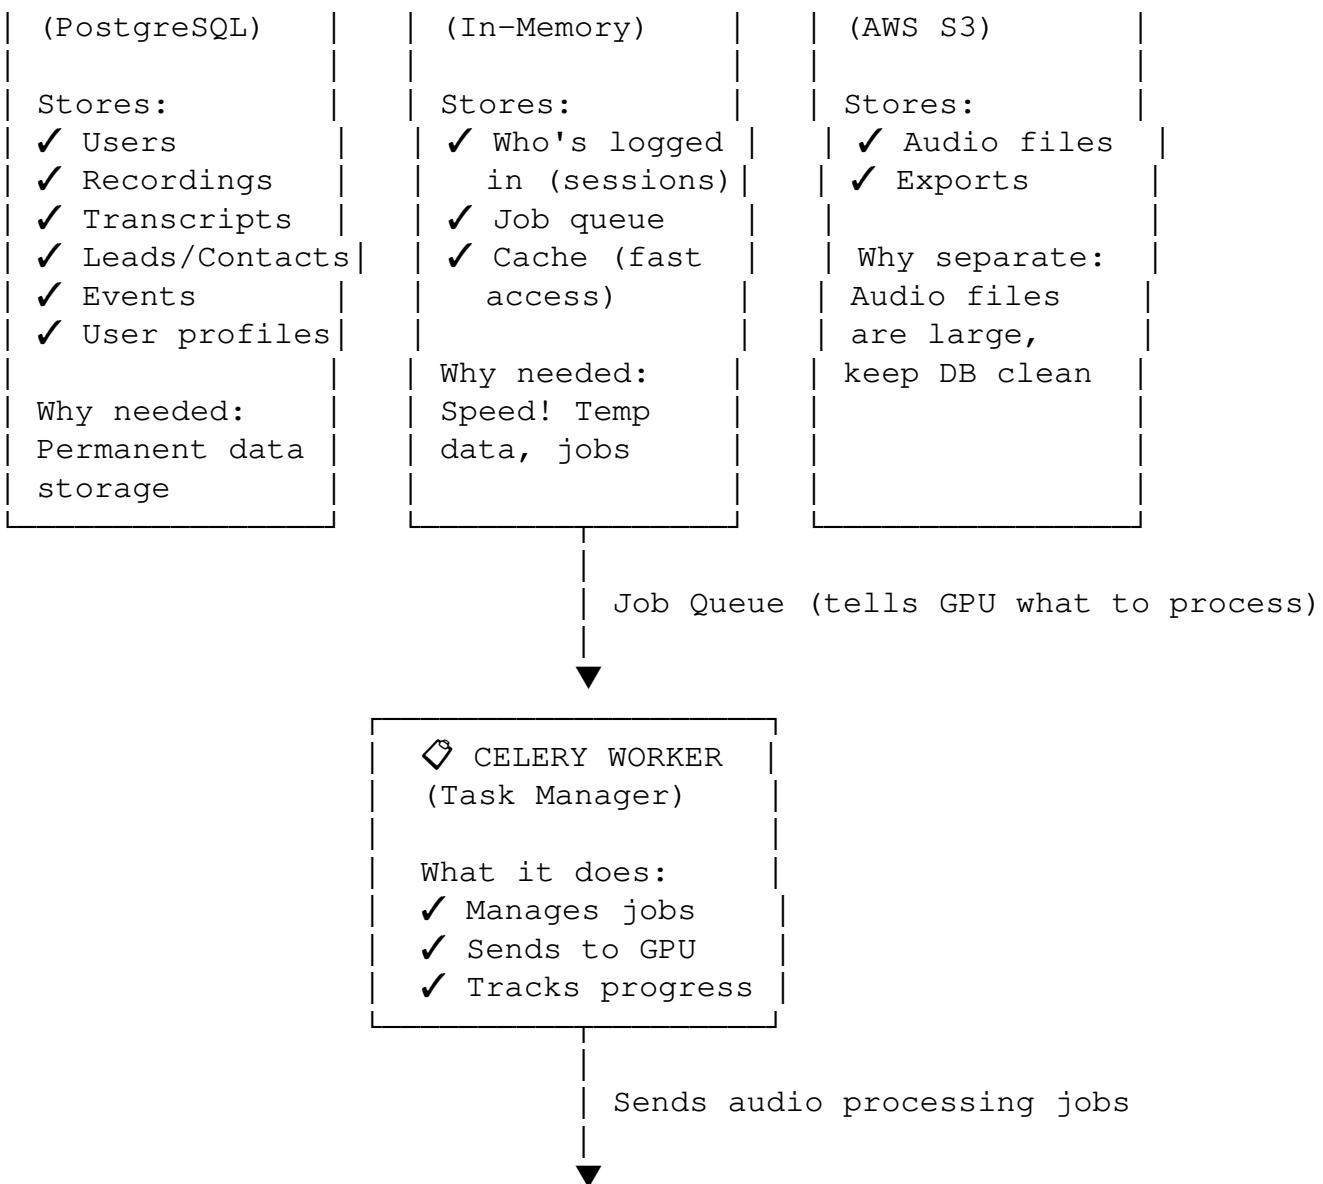


AYKA Platform - Complete Architecture (Non-Technical Friendly)

System Architecture Diagram





1 WHISPER AI (Speech to Text)
 - Converts audio → text
 - Processing time: 2–5 minutes for 30-min audio
 - Saves to: PostgreSQL (transcripts table)

2 PYANNOTE AI (Speaker Detection)
 - Identifies who said what
 - Labels: Speaker 1, Speaker 2, etc.
 - Saves to: PostgreSQL (speakers table)

3 LEAD GENERATION AGENT (CrewAI + GPT-4)
 - Reads transcript
 - Finds people mentioned

- Extracts: names, companies, roles, needs
- Searches LinkedIn for profiles
- Saves to: PostgreSQL (leads table)

- 4 EVENT DISCOVERY AGENT (CrewAI + GPT-4)
- Reads user profile (interests, location)
 - Searches Google for events
 - Filters by relevance
 - Ranks: must attend, should attend, nice to attend
 - Saves to: PostgreSQL (events table)

- 5 EMAIL AGENT (CrewAI + GPT-4)
- Creates beautiful HTML email
 - Includes leads + events
 - Sends via Gmail SMTP
 - Saves log to: PostgreSQL (email_logs table)

Total Processing Time: 3-8 minutes per recording

Why Each Component? (Simple Explanations)

1. Why 2 Servers (FastAPI on EC2)?

- **Reason:** If one server crashes, the other keeps working
- **Benefit:** Website stays online 24/7
- **Analogy:** Like having 2 cashiers at a store - if one is busy, the other serves customers

2. Why PostgreSQL Database?

- **Reason:** Store all user data permanently
- **What's stored:** Users, recordings, leads, events, everything
- **Analogy:** Like a filing cabinet that never loses data

3. Why Redis?

- **Reason 1:** Super fast temporary storage (100x faster than PostgreSQL)
- **Reason 2:** Manages job queue (who's next in line for processing)
- **Reason 3:** Remembers who's logged in
- **Analogy:** Like sticky notes for quick reminders vs. filing cabinet for permanent records

4. Why S3 Storage?

- **Reason:** Audio files are HUGE (30-min recording = 50-100MB)
- **Benefit:** Don't clog the database with large files

- **Analogy:** Like a warehouse for bulky items vs. office desk for documents

5. Why Separate GPU Server?

- **Reason:** AI processing needs powerful hardware (expensive)
- **Benefit:** Only pay for GPU when processing, not 24/7
- **Cost:** GPU = \$200/month, Regular server = \$30/month
- **Analogy:** Like renting a bulldozer only when you need it vs. owning a car

6. Why Load Balancer?

- **Reason:** Distributes users across servers evenly
 - **Benefit:** No single server gets overloaded
 - **Analogy:** Like a traffic cop directing cars to less busy lanes
-

Complete Data Flow (Step by Step)

1. User opens app/website
↓
2. User logs in
 - Server checks PostgreSQL (users table)
 - Redis stores session (you're logged in)
 ↓
3. User clicks "Record" → Audio captured
↓
4. Audio uploaded
 - Sent to S3 (audio storage)
 - Metadata saved in PostgreSQL (recordings table)
 - Job created in Redis queue
 ↓
5. GPU server picks up job
↓
6. Whisper converts speech → text
 - Saved to PostgreSQL (transcripts table)
 ↓
7. Pyannote identifies speakers
 - Saved to PostgreSQL (speakers table)
 ↓
8. Lead Agent analyzes transcript
 - Finds people, companies, opportunities
 - Searches LinkedIn
 - Saved to PostgreSQL (leads table)
 ↓
9. Event Agent searches for events
 - Finds relevant conferences/meetups
 - Ranks by relevance
 - Saved to PostgreSQL (events table)
 ↓
10. Email Agent sends summary
 - HTML email with leads + events
 - Saved to PostgreSQL (email_logs table)

- ↓
11. ✓ User sees results
 - API fetches from PostgreSQL
 - Redis caches for fast access
 - Displayed on app/website
-

Optional: Graph Database (Phase 2)

Neo4j - Why Later?

What it does: Connects people by shared interests

Example:

- You like: AI, Crypto, Startups
- John likes: AI, Crypto
- Sarah likes: AI, Startups
- Neo4j finds: "You should meet John and Sarah - you have common interests!"

Why not now: Adds complexity, not needed for MVP

When to add: When you have 1000+ users and want smart matching

Technology Choices - Why These?

Backend: FastAPI (Python)

- **Why:** Fast, modern, easy to add AI
- **Alternative:** Node.js/Express (slower for AI tasks)

Frontend: Next.js

- **Why:** Works on all devices, SEO-friendly
- **Alternative:** Plain React (no SEO)

Mobile: React Native

- **Why:** Write once, works on iOS + Android
- **Alternative:** Native apps (need 2 codebases)

Database: PostgreSQL

- **Why:** Reliable, handles millions of records, free
- **Alternative:** MongoDB (less structure, not ideal for our use case)

GPU: E2E Cloud

- **Why:** Cheaper than AWS GPU (50% less)
 - **Alternative:** AWS GPU (2x cost)
-

Cost Breakdown (Monthly)

Component	Why Needed	Cost
2 × API Servers (EC2)	Handle user requests, always online	\$60
PostgreSQL Database	Store all data permanently	\$20
Redis Cache	Fast access, job queue	\$15
S3 Storage	Store audio files	\$10
GPU Server (E2E)	AI processing (Whisper, Agents)	\$200
Load Balancer	Distribute traffic	\$20
TOTAL		\$325/month

For comparison:

- Hiring 1 person to manually process recordings: \$3000+ /month
 - Our AI does it automatically: \$325/month
-

Scalability - What Happens When You Grow?

10 Users

- 1 API server
- Small database
- Cost: \$150/month

100 Users

- 2 API servers
- Medium database
- Cost: \$325/month (current plan)

1,000 Users

- 5 API servers (auto-scaling)
- 2 GPU servers
- Larger database
- Cost: \$800/month

10,000 Users

- 20+ API servers
 - 5+ GPU servers
 - Database replicas
 - Cost: \$3,000-5,000/month
 - Revenue needed: \$10,000+ /month (break even at \$10/user)
-

Security & Privacy

How Data is Protected

1. **Encryption in Transit:** All data encrypted during transfer (HTTPS)
2. **Encryption at Rest:** Database encrypted on disk
3. **Access Control:** Only you see your data
4. **Backups:** Daily automatic backups (7 days)
5. **Audio Storage:** Can be auto-deleted after processing (GDPR compliant)

Who Can Access What?

- **Users:** Only their own data
 - **Admins:** Only with permission (audit logged)
 - **AI Agents:** No permanent storage, process and delete
-

Summary - Why This Architecture?

- ✓ **Reliable:** If one server fails, others keep working
- ✓ **Scalable:** Add more servers as users grow
- ✓ **Maintainable:** Modular design, easy to update
- ✓ **Fast:** Redis caching, optimized queries
- ✓ **Cost-effective:** Only pay for GPU when processing
- ✓ **Secure:** Industry-standard encryption and access control

Bottom Line: Professional, production-ready system that can handle growth from 10 to 10,000 users.